

CPSC 532W Assignment 4

Ali Seyfi - 97446637

Here is the link to the repository:

https://github.com/aliseyfi75/Probabilistic-Programming/tree/master/Assignment_4

write a black-box variational inference (VI) evaluator following section 4.4 of the book.

1 Code

Show code snippets that demonstrate the completeness and correctness of your BBVI implementation.

1.1 primitives

```
1 baseprimitives = {
2     '+' : lambda x: x[0] + x[1],
3     '-' : lambda x: x[0] - x[1],
4     '*' : lambda x: x[0] * x[1],
5     '/' : lambda x: x[0] / x[1],
6     '>' : lambda x: x[0] > x[1],
7     '>=' : lambda x: x[0] >= x[1],
8     '<' : lambda x: x[0] < x[1],
9     '<=' : lambda x: x[0] <= x[1],
10    '=' : lambda x: x[0] == x[1],
11    '=' : lambda x: torch.tensor([1]) if x[0] == x[1] else torch.tensor([0]),
12    'and' : lambda x: x[0] and x[1],
13    'or' : lambda x: x[0] or x[1],
14    'sqrt' : lambda x: torch.sqrt(x[0]),
15    'exp' : lambda x: torch.exp(x[0]),
16    'log' : lambda x: torch.log(x[0]),
17    'abs' : lambda x: torch.abs(x[0]),
18    'vector' : vector,
19    'list' : list,
20    'get' : get,
21    'put' : put,
22    'hash-map' : hash_map,
23    'first' : lambda x: x[0][0],
24    'last' : lambda x: x[0][-1],
25    'nth' : lambda x: x[0][int(x[1].item())],
26    'second' : lambda x: x[0][1],
27    'rest' : lambda x: x[0][1:],
28    'append' : append,
29    'cons' : lambda x: append([x[1], x[0]]),
30    'conj' : append,
31    'mat-add' : lambda x: x[0] + x[1],
32    'mat-mul' : lambda x: torch.matmul(x[0], x[1]),
33    'mat-transpose' : lambda x: x[0].T,
34    'mat-tanh' : lambda x: x[0].tanh(),
35    'mat-repmat' : lambda x: x[0].repeat((int(x[1].item()), int(x[2].item())))
36 }
```

Listing 1: primitives.py - Base primitives

```

1 def vector(x):
2     try:
3         vector = torch.stack(x)
4     except:
5         vector = x
6     return vector
7
8 def list(x):
9     try:
10        list = torch.stack(x)
11    except:
12        list = x
13    return list
14
15 def get(x):
16     if type(x[0]) == dict:
17         value = x[0][x[1].item()]
18     else:
19         value = x[0][x[1].long()]
20     return value
21
22 def put(x):
23     if type(x[0]) == dict:
24         x[0][x[1].item()] = x[2]
25     else:
26         x[0][x[1].long()] = x[2]
27     return x[0]
28
29 def hash_map(x):
30     keys = x[:, :2]
31     value = x[:, 2:]
32     new_keys = []
33     for key in keys:
34         try:
35             new_keys.append(key.item())
36         except:
37             new_keys.append(key)
38     result = dict(zip(new_keys, value))
39     return result
40
41 def append(x):
42     first = x[0]
43     second = x[1]
44
45     if first == 'vector':
46         first = torch.tensor([])
47     elif first.dim() == 0:
48         first = first.unsqueeze(0)
49     if second == 'vector':
50         second = torch.tensor([])
51     if second.dim() == 0:
52         second = second.unsqueeze(0)
53     return torch.cat((first, second))

```

Listing 2: primitives.py - Functions

```

1 import distributions as dist
2
3 class Dist:
4     def __init__(self, name, distribution, num_par, *par):
5         self.name = name
6         self.distribution = distribution
7         self.num_par = num_par
8         self.pars = []
9         for i in range(num_par):
10             self.pars.append(par[i])
11
12     def sample(self):
13         return self.distribution.sample()
14
15     def log_prob(self, c):
16         return self.distribution.log_prob(c)
17
18     def parameters(self):
19         return self.distribution.Parameters()
20
21     def make_copy_with_grads(self):
22         temp_dist = self.distribution
23         self.distribution = None
24         dist_copy = copy.deepcopy(self)
25         self.distribution = temp_dist
26         dist_copy.distribution = temp_dist.make_copy_with_grads()
27         return dist_copy
28
29 class normal(Dist):
30     def __init__(self, pars):
31         mean = pars[0]
32         var = pars[1]
33         normal_dist = dist.Normal(mean, var)
34         super().__init__('normal', normal_dist, 2, mean, var)
35
36 class beta(Dist):
37     def __init__(self, pars):
38         alpha = pars[0]
39         betta = pars[1]
40         super().__init__('beta', distributions.Beta(alpha, betta), 2, alpha, betta)
41
42 class exponential(Dist):
43     def __init__(self, par):
44         lamda = par[0]
45         super().__init__('exponential', distributions.Exponential(lamda), 1, lamda)
46
47 class uniform(Dist):
48     def __init__(self, pars):
49         a, b = pars[0], pars[1]
50         if a > b:
51             b = 5
52         uniform_dist = dist.Uniform(a, b)
53         super().__init__('uniform', uniform_dist, 2, a, b)
54
55 class discrete(Dist):
56     def __init__(self, pars):
57         prob = pars[0]
58         discrete_dist = dist.Categorical(prob)
59         super().__init__('discrete', discrete_dist, 0)
60
61 class bernoulli(Dist):
62     def __init__(self, pars):
63         p = pars[0]
64         bernoulli_dist = dist.Bernoulli(p)

```

```

65         super().__init__('bernoulli', bernoulli_dist, 1, p)
66
67 class gamma(Dist):
68     def __init__(self, pars):
69         alpha, betta = pars[0], pars[1]
70         gamma_dist = dist.Gamma(alpha, betta)
71         super().__init__('gamma', gamma_dist, 2, alpha, betta)
72
73 class dirichlet(Dist):
74     def __init__(self, pars):
75         dirichlet_dist = dist.Dirichlet(pars[0])
76         super().__init__('dirichlet', dirichlet_dist, len(pars), *pars)
77
78 class dirac(Dist):
79     def __init__(self, value):
80         mean = value[0]
81         mean = torch.clip(mean, -1e5, 1e5)
82         var = torch.tensor(1e-5)
83         super().__init__('normal', distributions.Normal(mean, var), 2, mean, var)

```

Listing 3: primitives.py - Distributions

```

1 distlist = {
2     'normal' : normal,
3     'beta' : beta,
4     'exponential' : exponential,
5     'uniform' : uniform,
6     'discrete' : discrete,
7     'bernoulli': bernoulli,
8     'gamma': gamma,
9     'dirichlet': dirichlet,
10    'flip': bernoulli,
11    'dirac': dirac,
12    'uniform-continuous': uniform
13 }

```

Listing 4: primitives.py - distlist

1.2 evaluation based sampling

This part is same as last assignment.

1.3 topological sort

This part is same as last assignment.

1.4 BBVI evaluator

```
1 def BBVI_evaluator(order_node, graph, sigma):
2     P = graph[1]['P']
3     Y = graph[1]['Y']
4     Q = sigma['Q']
5     G = sigma['G']
6     optimizer = sigma['optimizer']
7     results = {}
8
9     for node in order_node:
10         link_function = P[node][0]
11
12         if link_function == 'sample*':
13             d = deterministic_eval(value_subs(P[node][1], results))
14             if node not in Q:
15                 Q[node] = d.make_copy_with_grads()
16                 optimizer[node] = torch.optim.Adam(Q[node].parameters(), lr=0.01)
17             result = Q[node].sample()
18             G[node] = grad_log_prob(Q[node], result)
19             try:
20                 sigma_temp = d.log_prob(result) - Q[node].log_prob(result)
21                 sigma['logW'] += sigma_temp
22             except:
23                 sigma['logW'] += 0
24
25         elif link_function == 'observe*':
26             result = torch.tensor(Y[node])
27             d = deterministic_eval(value_subs(P[node][1], results))
28             sigma_temp = d.log_prob(result)
29             sigma['logW'] += sigma_temp
30
31         results[node] = result
32
33     return results, sigma
```

Listing 5: graph_based_sampling.py - BBVIevaluator

1.5 grad log prob

```
1 def grad_log_prob(dist, value):
2     for param in dist.parameters():
3         param = param.clone().detach()
4         param.requires_grad = True
5     log_prob = dist.log_prob(value)
6     log_prob.backward()
7     grad = [param.grad for param in dist.parameters()]
8     return grad
```

Listing 6: graph_based_sampling.py - grad_log_prob

1.6 BBVI

```
1 def BBVI(graph, T, L):
2     sigma = {'Q': {}, 'optimizer': {}}
3     order_node = topological_sort(graph)
4
5     results = []
6     log_weights = []
7     posteriors = []
8
9     for t in range(T):
10         sigma['G'] = {}
11         gradients = []
12         log_ws = []
13
14         for l in range(L):
15             sigma['logW'] = 0
16             result, sigma = BBVI_evaluator(order_node, graph, sigma)
17             gradients.append(copy.deepcopy(sigma['G']))
18             log_ws.append(sigma['logW'])
19
20         if t==0:
21             posteriors.append(copy.deepcopy(sigma['Q']['sample2'].parameters()))
22
23         ELBO_gradients(gradients, log_ws, sigma['Q'])
24
25         for optimizer in sigma['optimizer'].values():
26             optimizer.step()
27             optimizer.zero_grad()
28
29         post_temp = {}
30         for q in sigma['Q']:
31             post_temp[q] = sigma['Q'][q].parameters().copy()
32
33         posteriors.append(post_temp)
34         result_temp = deterministic_eval(value_subs(graph[2], result))
35         results.append(result_temp)
36         log_weights.append(log_ws[-1])
37         wandb.log({'ELBO': torch.mean(torch.stack(log_weights)).detach().numpy()})
38
39     return results, log_weights, posteriors
```

Listing 7: graph_based_sampling.py - BBVI

1.7 infinity skipper

```
1 def inf_skipper(gradients, log_ws):
2     temp_gradients = []
3     temp_log_ws = []
4
5     for i in range(len(log_ws)):
6         if log_ws[i] == float('-inf'):
7             continue
8         temp_gradients.append(gradients[i])
9         temp_log_ws.append(log_ws[i])
10
11     return temp_gradients, temp_log_ws
```

Listing 8: graph_based_sampling.py - inf_skipper

1.8 ELBO gradient

```
1 def ELBO_gradients(gradients, log_ws, posteriors):
2
3     gradients, log_ws = inf_skipper(gradients, log_ws)
4     len_grads = len(gradients)
5
6     var_union = list(set([var for grad in gradients for var in grad]))
7
8     Fs = []
9     Gs = []
10    stack = {}
11
12    for var in var_union:
13        gradient_var = gradients[0][var]
14        if len(gradient_var[0].shape) > 0 and len(gradient_var[0]) > 1:
15            gradient_var = [grad.clone().detach().requires_grad_(True) for grad in
16            gradient_var[0]]
17            stack[var] = len(gradient_var)
18
19            len_vars = len(gradient_var)
20
21            G_var = torch.zeros((len_grads, len_vars))
22            F_var = torch.zeros((len_grads, len_vars))
23
24            for lg in range(len_grads):
25                G_var[lg, :] = torch.stack(gradients[lg][var])
26                F_var[lg, :] = G_var[lg, :] * log_ws[lg]
27            Gs.append(G_var.detach().numpy())
28            Fs.append(F_var.detach().numpy())
29
30    Gs = np.column_stack(Gs)
31    Fs = np.column_stack(Fs)
32
33    num = np.sum([np.cov(Fs[:, v], Gs[:, v])[0, 1] for v in range(Gs.shape[1])])
34    denum = np.sum([np.var(Gs[:, v]) for v in range(Gs.shape[1])])
35    b_hat = 0.
36    if not denum == 0. and not np.isnan(num):
37        b_hat = num/denum
38
39    counter_1 = 0
40    for var in var_union:
41        gradient_var = gradients[0][var]
42        counter_2 = len(gradient_var)
43        if var in stack:
44            counter_2 = stack[var]
45        g_hat = np.array([np.sum(Fs[:, v] - b_hat * Gs[:, v]) / len_grads for v in range(
46        counter_1, counter_1+counter_2)])
47        if var in stack:
48            g_hat = [g_hat]
49        for i, parameter in enumerate(posteriors[var].parameters()):
50            parameter.grad = torch.tensor(-g_hat[i], dtype=parameter.grad.dtype)
51        counter_1 += counter_2
52    return
```

Listing 9: graph_based_sampling.py - ELBO_gradient

2 Results

2.1 Task 1

$T = 10^4$ and $L = 50$ for this task.

Time of drawing samples: **410.81 seconds**

Posterior expected value of mu is: **7.3007**

Parameters of the posterior distribution of mu: $\mu = 7.2742$ and $\sigma = 0.4931$.




Figure 1: Histogram of posterior distribution of mu




Figure 2: ELBO in task 1

2.2 Task 2

$T = 5 * 10^3$ and $L = 50$ for this task.

Time of drawing samples: **510.36 seconds**

Posterior mean of slope is: **2.1169**

Posterior mean of bias is: **-0.4039**




Figure 3: ELBO in task 2

2.3 Task 3

$T = 2 * 10^3$ and $L = 50$ for this task.

Time of drawing samples: **581.53 seconds**

Posterior mean of probability that the first and second datapoint are in the same cluster is: **0.6715**

Here we have symmetric between our states due to our choice of prior, so our optimization model will converge to different solutions each time we run the program, which is actually the mode-seeking behaviour. Also during the optimization, we can see jumps in ELBO values, and the reason of that is actually jumping between modes and show the mode-seeking behaviour of this method.

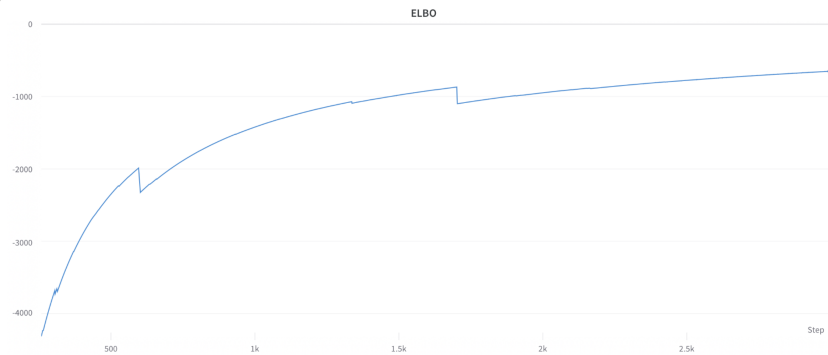


Figure 4: ELBO in task 3

2.4 Task 4

Time of drawing samples: **595.87 seconds**




Figure 5: Posterior distribution of b_0




Figure 6: Posterior distribution of w_0




Figure 7: Posterior distribution of b_1




Figure 8: Posterior distribution of w_1



Figure 9: ELBO in task 4

In BBVI method, we need proposal distributions to be differentiable. However, in parameter estimation via gradient descent we need both proposal distribution and joint distribution to be differentiable. So the advantage of BBVI will be working with discrete and continuous models, but the variance could be a problem in this case. But in the other method we have generally a better behaved variance.

2.5 Task 5

$T = 2 * 10^3$ and $L = 25$ for this task.

Time of drawing samples: **45.19 seconds**

learned variational distribution for s: **Uniform(0.9024, 1.7703)**

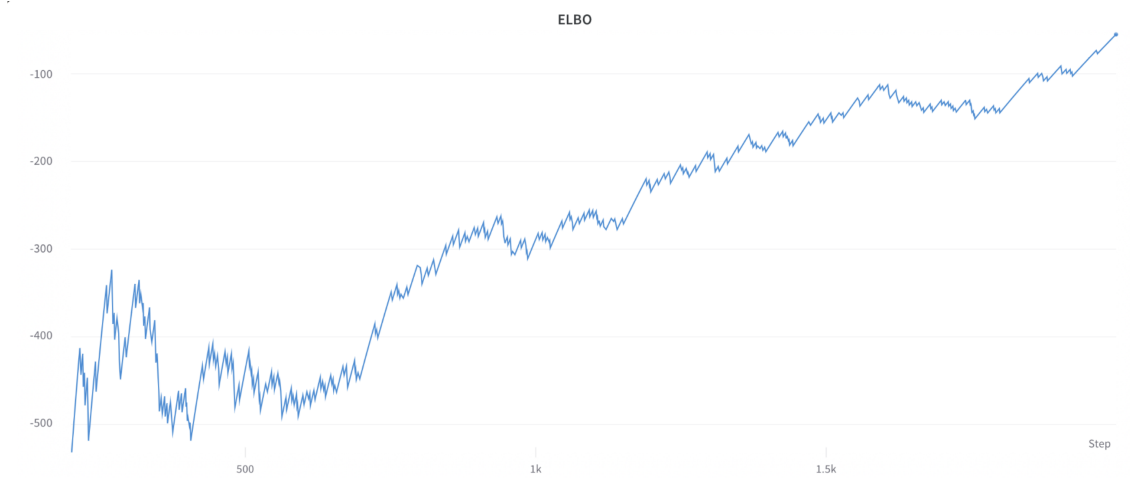


Figure 10: ELBO in task 5